

Baseline evaluation scorecard

Held-out eval uses greedy preset and 64 test batches. Accuracy values are percentages.

Base model

51.56%

33/64 correct

Best LoRA ckpt

28.13%

step 2000, 18/64 correct

Final LoRA

3.13%

step 3364, 2/64 correct

Main result

- The final trained LoRA checkpoint is worse than the frozen base model and worse than the best early LoRA checkpoint.
- Checkpoint-wise eval shows degradation after step 2000: 28.13% -> 20.31% -> 6.25% -> 3.13%.
- The report treats step 2000 as the best observed LoRA checkpoint and step 3364 as evidence of late-stage collapse.