


Reward and KL timeline

Reward should improve while KL stays controlled; this run improves early then collapses late.



The best checkpoint by held-out accuracy is step 2000; final step 3364 is substantially worse.
Late eval reward is negative while response failure rates rise, consistent with training collapse rather than successful alignment.