

Train reward by aligned rollouts

Report-facing reward values are rewritten to the shared baseline 0-10 scale; native values are retained as value_native.

