

Training reward scaled to native maximum

Each total reward is divided by that reward mechanism's theoretical maximum: baseline=10.0, simple=1.2, format-aware=1.8.

